

Lecture 10

One-hot encoding and Word Count

Course Outcome:

1. Apply One-Hot Encoding and Bag of Words to a small dataset.
2. Analyze the Strengths and Limitations of OHE & BoW

Introduction- One-hot encoding and Word Count

Both One-Hot Encoding (OHE) and Word Count (Bag of Words - BOW) are fundamental text representation techniques in Natural Language Processing (NLP).

One-hot encoding

One-Hot Encoding is a technique used in Natural Language Processing (NLP) and Machine Learning to represent categorical data as binary vectors.

 Key Idea:

- Each unique word (or category) is represented as a vector of 0s and 1s.
- The word's position in the vocabulary (vocab) is marked as 1, while all other positions remain 0.
- Used in text processing, sentiment analysis, deep learning (Neural Networks, RNNs, Transformers).

Question 1. Given a set of sentences, how do we represent words numerically so that a computer can process them?

Think: What if we assign each word a unique identifier?

Ans: step 1: Create the Vocabulary of Unique Words List:

Sample Dataset

Let's assume we have the following sentences in our dataset:

1. "The phone is excellent."
2. "The battery performance is excellent."
3. "The phone has excellent battery life."

Index	Word	One-Hot Encoding
0	The	[1, 0, 0, 0, 0, 0, 0, 0]
1	phone	[0, 1, 0, 0, 0, 0, 0, 0]
2	is	[0, 0, 1, 0, 0, 0, 0, 0]
3	excellent	[0, 0, 0, 1, 0, 0, 0, 0]
4	battery	[0, 0, 0, 0, 1, 0, 0, 0]
5	performance	[0, 0, 0, 0, 0, 1, 0, 0]
6	has	[0, 0, 0, 0, 0, 0, 1, 0]
7	life	[0, 0, 0, 0, 0, 0, 0, 1]

Step 1: Create the Vocabulary

Extract unique words across all text samples

Vocab: ["The", "phone", "is", "excellent", "battery", "performance", "has", "life"]

Question 2 : How can we represent the first review "The phone is excellent"

Answer:

📌 Each word is assigned a unique vector with 1 where the word appears and 0 elsewhere.

Sentence 2: The battery performance is excellent.

```
[  
  [1, 0, 0, 0, 0, 0, 0, 0],  
  [0, 0, 0, 0, 1, 0, 0, 0],  
  [0, 0, 0, 0, 0, 1, 0, 0],  
  [0, 0, 1, 0, 0, 0, 0, 0],  
  [0, 0, 0, 1, 0, 0, 0, 0]  
]
```

Sentence 3: The phone has excellent battery life.

```
[  
  [1, 0, 0, 0, 0, 0, 0, 0],  
  [0, 1, 0, 0, 0, 0, 0, 0],  
  [0, 0, 0, 0, 0, 0, 1, 0],  
  [0, 0, 0, 1, 0, 0, 0, 0],  
  [0, 0, 0, 0, 1, 0, 0, 0],  
  [0, 0, 0, 0, 0, 0, 0, 1]  
]
```

Question 2: Can you come up with a different Approach?

◆ Example: Sentence-Level One-Hot Encoding

📌 Sentences: 1 "The phone is excellent."

2 "The battery performance is excellent."

3 "The phone has excellent battery life."

✓ Step 1: Create a Vocabulary

👉 Extract **unique words** from all sentences:

📌 Vocab: ["The", "phone", "is", "excellent", "battery", "performance", "has", "life"]

✓ Step 2: Convert Sentences into One-Hot Encoding Vectors

Sentence	The	phone	is	excellent	battery	performance	has	life
"The phone is excellent."	1	1	1	1	0	0	0	0
"The battery performance is excellent."	1	0	1	1	1	1	0	0
"The phone has excellent battery life."	1	1	0	1	1	0	1	1

Analogy for One-Hot Encoding: The Word Pool and Vocabulary Selection

◆ Imagine a Pool of Words (Vocabulary)

Think of all the words in a language as a huge pool of water.

- This pool contains millions of words—both useful and irrelevant for our task.
- But we don't need all words, so we create a small vocabulary (subset of words) that is important for our task.

✓ Example: If we want to analyze sentiment in restaurant reviews, we only select words like: ["delicious", "amazing", "terrible", "bad", "service", "slow"]

✗ We ignore words like "chair", "table", "fork" because they don't affect sentiment.

◆ One-Hot Encoding: The Selected Words Become Unique Identifiers

📌 Now, each word in our vocabulary gets a unique identity (vector)—like assigning it a seat in a stadium.

✅ Example: Suppose our vocabulary is:
👉 ["delicious", "amazing", "bad", "terrible", "slow", "great"]

Now, we encode each word as a binary vector where:

- 1 means the word is present.
- 0 means it is absent.

✅ **Sentence-level One-Hot Encoding:** Used for tasks like **text classification** and **sentiment analysis** (treats sentence as a whole).

✅ **Word-level One-Hot Encoding:** Used for tasks like **word embeddings**, **RNNs**, and **sequence modeling** (treats words individually).

Limitations of one-hot encoding?

🔴 **Limitations: Why One-Hot Encoding is Being Replaced?**

Problem	Why It's an Issue?	Better Alternative
High Dimensionality	If we have 10,000 words, we need a 10,000-column matrix	✅ Word Embeddings (Word2Vec, BERT)
Doesn't Capture Meaning	"Great" and "Awesome" are treated as separate, unrelated words	✅ Embeddings (Vectors capture meaning)
Memory Inefficiency	Mostly sparse (many 0s), wasting memory	✅ Compressed Embedding Models

1. **High Dimensionality (Curse of Dimensionality):** A vocabulary of 10,000 words requires each word to be represented by a 10,000-dimensional vector, with only one active element (1), leading to wasted memory.
2. **Lack of Semantic Meaning:** One-hot encoding treats all words as independent entities without any notion of similarity. Example: The words "**King**" and "**Queen**" would have completely different one-hot vectors, even though they are semantically related.
3. **Not Scalable for Large Datasets:** In real-world NLP tasks (e.g., training GPT models), vocabularies can reach millions of words, making one-hot encoding impractical.
4. **No Context Awareness:** One-hot encoding does not consider word meaning or word usage in different contexts.

Example:

Sentence 1: "He went to the bank to withdraw money."

Sentence 2: "She sat by the bank of the river."

The word "bank" has different meanings but is treated identically in one-hot encoding.

Applications of One hot encoding:

Used in Text Processing before Moving to Advanced Embeddings

- Before deep learning, One-Hot Encoding and word count was a standard way to represent words for machine learning models.
- Many datasets contain categorical features (e.g., "Yes/No", "Red/Blue/Green"). One-Hot Encoding converts these categories into numbers so that machine learning models can process them.

Code of OHE:

```
from sklearn.preprocessing import OneHotEncoder
import numpy as np

# Given words in reviews
vocab = np.array(["the", "phone", "is", "great", "battery", "bad", "has", "good", "life"]).reshape(-1, 1)

# Initialize One-Hot Encoder
encoder = OneHotEncoder(sparse=False)

# Apply One-Hot Encoding
one_hot_encoded = encoder.fit_transform(vocab)

# Display Encoded Vectors
for word, encoding in zip(vocab.flatten(), one_hot_encoded):
    print(f'{word}: {encoding}')
```

Word Count

Word count is a basic text processing technique used to determine **how many times each word appears in a given document or text corpus**.

1 Basic Word Count Example

Let's take a simple sentence:

🗨️ **Sentence:**

📌 *"The food was good, and the service was good too."*

◆ Step 1: Tokenize (Split) Words

We remove punctuation and split the sentence into words:

```
["the", "food", "was", "good", "and", "the", "service", "was", "good", "too"]
```

◆ Step 2: Count Word Occurrences

Word	Count
the	2
food	1
was	2
good	2
and	1
service	1
too	1

✅ Final Word Count Representation:

```
{"the": 2, "food": 1, "was": 2, "good": 2, "and": 1, "service": 1, "too": 1}
```

Review is represented as:

```
[2,1,2,2,1,2,1,2,1]
```

Applications of Bag of Words (BOW):

1. Text Classification & Sentiment Analysis

✅ **BoW is used to classify text into categories** such as spam detection, fake news detection, or sentiment analysis.

💡 **Example:**

- ◆ **Task:** Classify movie reviews as **positive** or **negative**.
- ◆ **Dataset:**

- **Review 1:** "The movie was excellent, I loved it!" ✅ (Positive)
- **Review 2:** "Worst movie ever, complete waste of time!" ❌ (Negative)

📌 How BoW Helps?

- Words like "excellent", "loved" appear more in **positive** reviews.
- Words like "worst", "waste" appear more in **negative** reviews.
- **ML models (Naïve Bayes, Logistic Regression, SVM)** use these BoW features to classify reviews.

2. Word Cloud Example Using Bag of Words (BoW) in Python

🚀 What is a Word Cloud?

A **Word Cloud** is a **visual representation of text data**, where the size of each word represents its **frequency or importance**. It is still widely used in:

- ✅ **Customer Feedback Analysis** (Amazon, Flipkart, Google Reviews)
- ✅ **Social Media Analysis** (Twitter trends, YouTube comments)
- ✅ **News Analytics** (Finding trending topics in news articles)
- ✅ **SEO & Content Marketing** (Finding frequently used keywords)

Limitations of BOW?

❌ 1. Ignores Word Order (Loses Context)

📌 Problem:

BoW does not consider **word sequence**, so "not happy" and "happy not" are treated **the same**, even though they have different meanings.

💡 Example:

- **Sentence 1:** "The movie was not good." ❌ (Negative)
- **Sentence 2:** "The movie was good." ✅ (Positive)

🚀 Why is this bad?

- Both sentences may have **similar BoW word counts** ("movie", "was", "good"), but the **first one is negative while the second is positive**.

✅ **Solution:** Use **N-Grams (bigrams, trigrams)** or **Transformers** to preserve word order.

✗ 2. High Dimensionality & Sparsity

📌 Problem:

BoW creates **huge sparse matrices** because each document is represented as a **vector with as many features as the vocabulary size**.

💡 Example:

- If the vocabulary size = **100,000 words**, then each document will be a **100,000-dimensional vector**, mostly filled with **zeros**.

🚀 Why is this bad?

- Consumes a lot of memory.
- Slows down training and prediction.

✅ **Solution:** Use **TF-IDF** to reduce feature space or **dimensionality reduction techniques** like **PCA**.

✗ 3. Cannot Capture Word Meaning or Synonyms

📌 Problem:

BoW **treats synonyms as separate features**, ignoring the fact that **"amazing"**, **"great"**, and **"fantastic"** all have similar meanings.

💡 Example:

- **"The food was amazing!"** → **["food", "was", "amazing"]**
- **"The food was great!"** → **["food", "was", "great"]**

🚀 Why is this bad?

- **"amazing"** and **"great"** should contribute **similarly to sentiment analysis**, but BoW treats them as completely **different words**.

✅ **Solution:** Use **Word Embeddings (Word2Vec, GloVe, BERT)** that learn word relationships.

✗ 4. Fails for Long Documents & Real-Time Applications

📌 Problem:

- Longer documents naturally have **higher word counts**, leading to **biased models**.

- BoW models need **retraining every time new data is added**, making it **inefficient for real-time NLP**.

💡 **Example:**

- A **500-word blog post** will have much higher word counts than a **10-word customer review**, even if both talk about **the same product**.

🚀 **Why is this bad?**

- BoW models **favor longer documents** because they have **more words**, even if the sentiment or topic is the same.

✅ **Solution:**

- Use **TF-IDF (which normalizes word importance instead of raw count)**.
- Use **Neural Networks with Embeddings** instead of BoW for large-scale NLP.

📌 **Summary: Comparing OHE & BoW Limitations**

Limitation	One-Hot Encoding (OHE)	Bag of Words (BoW)
❌ High Dimensionality	Needs large binary vectors	Needs large word count vectors
❌ Ignores Word Order	No sequence information	No sequence information
❌ Does Not Capture Meaning	"great" and "excellent" treated separately	"happy" and "joyful" treated separately
❌ Fails with Out-of-Vocabulary Words	Cannot handle new words	Cannot handle new words
❌ Not Suitable for Large Datasets	Huge sparse matrices	High memory consumption
❌ Poor Performance on Long Texts	All words have equal weight	Longer texts dominate shorter ones

✅ **Better Alternatives:**

- **Word Embeddings (Word2Vec, FastText, BERT)** → Capture meaning & context.
- **TF-IDF** → Reduces high dimensionality & normalizes word importance.
- **Transformers (BERT, GPT-4)** → Handle word relationships, context, and meaning.

Code of BOW

<https://colab.research.google.com/drive/1puXZaq6QKGMylth3CqsDwZQcr7PMom26#scrollTo=3rv5J3eR4B3N>